

# FacialEmotionRecognition

June 2, 2022

```
[1]: # Basics
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Sklearn
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.decomposition import PCA
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.cluster import KMeans
from sklearn.pipeline import Pipeline
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# Tensorflow
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
import tensorflow.keras.layers as layers
from tensorflow.keras.utils import to_categorical

# Graphing Style
%matplotlib inline

[4]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

# 1 Data Cleaning and Exploratory Data Analysis

```
[5]: face = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/face.csv")
```

```
[6]: face.shape
```

```
[6]: (28709, 2)
```

```
[7]: face["pixels"] = face.pixels.apply(lambda x: np.array(tuple(map(int, x.  
    ↪split()))))
```

```
[8]: face.head()
```


```
[8]:      emotion                                pixels  
0         0  [70, 80, 82, 72, 58, 58, 60, 63, 54, 58, 60, 4...  
1         0  [151, 150, 147, 155, 148, 133, 111, 140, 170, ...  
2         2  [231, 212, 156, 164, 174, 138, 161, 173, 182, ...  
3         4  [24, 32, 36, 30, 32, 23, 19, 20, 30, 41, 21, 2...  
4         6  [4, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3, 15, 23...
```

```
[9]: face.emotion.value_counts()
```

```
[9]: 3    7215  
     6    4965  
     4    4830  
     2    4097  
     0    3995  
     5    3171  
     1     436  
     Name: emotion, dtype: int64
```

```
[10]: sns.histplot(face.emotion)
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e8769b4d0>
```



```
[11]: emo_di = {0: "Angry", 1: "Disgust", 2: "Fear", 3: "Happy", 4: "Sad", 5: "Surprise", 6: "Neutral"}
```

```
[12]: plt.figure(figsize = (20,20))
start_index = 0
for i in range(7):
    plt.subplot(1,7,i+1)
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])
    plt.imshow(face[face.emotion == i].pixels.iloc[20].reshape(48,48))
    plt.xlabel("{} = {}".format(i, emo_di[i]))
```



For each one, consider using Cross Validation and/or PCA

```
[13]: train_x, test_x, train_y, test_y = train_test_split(np.concatenate(np.
    ↳ asarray(face["pixels"])).reshape(-1, 48 * 48),
```

```

face.emotion,
train_size=0.7,
random_state=1)
val_x, test_x, val_y, test_y = train_test_split(test_x, test_y, test_size=0.5,
↪random_state=1)

x_mean = np.mean(train_x)
x_std = np.std(train_x) + 1e-10
train_x = (train_x - x_mean) / x_std
val_x = (val_x - x_mean) / x_std
test_x = (test_x - x_mean) / x_std

```

## 2 Modeling

### 2.1 Naive Bayes

```

[14]: nb = GaussianNB()
      nb.fit(train_x, train_y)

```

```

[14]: GaussianNB()

```

```

[15]: nb.score(train_x, train_y)

```

```

[15]: 0.2174562101910828

```

```

[16]: nb.score(test_x, test_y)

```

```


[16]: 0.2173206408172742

```

```

[19]: cm = confusion_matrix(test_y, nb.predict(test_x))
      cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
      fig, ax = plt.subplots(figsize=(10,10))
      sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=nb.classes_, yticklabels=nb.
↪classes_)
      plt.title('Confusion Matrix of Naive Bayes Classifier')
      plt.ylabel('Actual')
      plt.xlabel('Predicted')
      plt.show(block=False)

```



## 2.2 K-Means + Logistic Regression

```
[20]: pl = Pipeline([
    ("kmeans", KMeans(n_clusters=35)),
    ("logisreg", LogisticRegression(max_iter=10000))
])
```

```
[21]: pl.fit(train_x, train_y)
```

```
[21]: Pipeline(steps=[('kmeans', KMeans(n_clusters=35)),
    ('logisreg', LogisticRegression(max_iter=10000))])
```


```
[22]: pl.score(train_x, train_y)
```

```
[22]: 0.2910031847133758
```

```
[23]: pl.score(test_x, test_y)
```


```
[23]: 0.3018342233573253
```

```
[24]: cm = confusion_matrix(test_y, pl.predict(test_x))
      cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
      fig, ax = plt.subplots(figsize=(10,10))
      sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=nb.classes_, yticklabels=nb.
      ↪classes_)
      plt.title('Confusion Matrix of K-Means + Logistic Regression')
      plt.ylabel('Actual')
      plt.xlabel('Predicted')
      plt.show(block=False)
```



## 2.3 KNN with PCA + CV

```
[25]: pca = PCA().fit(train_x)
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('number of components')
plt.ylabel('cumulative explained variance');
```



```
[26]: pl = Pipeline([
        ("PCA", PCA(n_components=100)),
        ("knn", KNeighborsClassifier())
    ])
    params = {"knn__n_neighbors": [1, 3, 5, 7, 12, 15, 20]}
    grids = GridSearchCV(pl, params, cv=5)
    grids.fit(train_x, train_y)
```

```
[26]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('PCA', PCA(n_components=100)),
                                             ('knn', KNeighborsClassifier())]),
                  param_grid={'knn__n_neighbors': [1, 3, 5, 7, 12, 15, 20]})
```

```
[27]: grids.best_params_
```

```
[27]: {'knn__n_neighbors': 1}
```

```
[28]: grids.best_score_
```


```
[28]: 0.3485272195875561
```

```
[29]: grids.best_estimator_.score(test_x, test_y)
```

```
[29]: 0.3805433016020432
```



```
[30]: cm = confusion_matrix(test_y, grids.best_estimator_.predict(test_x))
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=nb.classes_, yticklabels=nb.
↪classes_)
plt.title('Confusion Matrix of KNN with PCA + CV')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show(block=False)
```



## 2.4 Random Forest

```
[31]: rfc = RandomForestClassifier(max_depth=5)
      rfc.fit(train_x, train_y)
```

```
[31]: RandomForestClassifier(max_depth=5)
```


```
[32]: rfc.score(train_x, train_y)
```

```
[32]: 0.3391222133757962
```

```
[33]: rfc.score(test_x, test_y)
```

```
[33]: 0.3262131413977246
```

```
[35]: cm = confusion_matrix(test_y, rfc.predict(test_x))
      cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
      fig, ax = plt.subplots(figsize=(10,10))
      sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=nb.classes_, yticklabels=nb.
        ↪classes_)
      plt.title('Confusion Matrix of Random Forest Classifier')
      plt.ylabel('Actual')
      plt.xlabel('Predicted')
      plt.show(block=False)
```



## 2.5 Support Vector Machine with PCA

```
[36]: pl = Pipeline([
        ("PCA", PCA(n_components=50)),
        ("svm", SVC())
    ])
    pl.fit(train_x, train_y)
```

```
[36]: Pipeline(steps=[('PCA', PCA(n_components=50)), ('svm', SVC())])
```


```
[37]: pl.score(train_x, train_y)
```

[37]: 0.528015525477707

```
[38]: pl.score(test_x, test_y)
```

[38]: 0.4211748316693754

```
[39]: cm = confusion_matrix(test_y, pl.predict(test_x))
      cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
      fig, ax = plt.subplots(figsize=(10,10))
      sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=nb.classes_, yticklabels=nb.
        ↪classes_)
      plt.title('Confusion Matrix of SVM with PCA')
      plt.ylabel('Actual')
      plt.xlabel('Predicted')
      plt.show(block=False)
```



## 2.6 Artificial Neural Network

```
[40]: train_y_cat = to_categorical(train_y)
      val_y_cat = to_categorical(val_y)
      test_y_cat = to_categorical(test_y)
```

```
[41]: model = Sequential()
      model.add(Dense(256, activation = 'relu', input_shape = (48 * 48,)))
      model.add(layers.Dropout(0.2))
      model.add(Dense(256, activation = 'relu'))
      model.add(layers.Dropout(0.5))
```

```

model.add(Dense(7, activation = 'softmax'))

model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = [
    ↪ 'accuracy'])
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	590080
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 256)	65792
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 7)	1799

=====  
 Total params: 657,671  
 Trainable params: 657,671  
 Non-trainable params: 0  
 =====

```

[42]: history = model.fit(train_x, train_y_cat, epochs = 20, validation_data = (
    ↪ val_x, val_y_cat))

```

```

Epoch 1/20
628/628 [=====] - 3s 3ms/step - loss: 1.9157 -
accuracy: 0.2620 - val_loss: 1.7005 - val_accuracy: 0.3281
Epoch 2/20
628/628 [=====] - 2s 3ms/step - loss: 1.7268 -
accuracy: 0.3102 - val_loss: 1.6687 - val_accuracy: 0.3416
Epoch 3/20
628/628 [=====] - 2s 3ms/step - loss: 1.6865 -
accuracy: 0.3293 - val_loss: 1.6284 - val_accuracy: 0.3595
Epoch 4/20
628/628 [=====] - 2s 3ms/step - loss: 1.6627 -
accuracy: 0.3350 - val_loss: 1.6245 - val_accuracy: 0.3500
Epoch 5/20
628/628 [=====] - 2s 3ms/step - loss: 1.6454 -
accuracy: 0.3452 - val_loss: 1.5907 - val_accuracy: 0.3739
Epoch 6/20
628/628 [=====] - 2s 3ms/step - loss: 1.6265 -
accuracy: 0.3529 - val_loss: 1.5924 - val_accuracy: 0.3688

```

```

Epoch 7/20
628/628 [=====] - 2s 3ms/step - loss: 1.6011 -
accuracy: 0.3652 - val_loss: 1.5835 - val_accuracy: 0.3716
Epoch 8/20
628/628 [=====] - 2s 3ms/step - loss: 1.5847 -
accuracy: 0.3729 - val_loss: 1.5692 - val_accuracy: 0.3908
Epoch 9/20
628/628 [=====] - 2s 3ms/step - loss: 1.5703 -
accuracy: 0.3796 - val_loss: 1.5744 - val_accuracy: 0.3804
Epoch 10/20
628/628 [=====] - 2s 3ms/step - loss: 1.5537 -
accuracy: 0.3880 - val_loss: 1.5695 - val_accuracy: 0.3864
Epoch 11/20
628/628 [=====] - 2s 3ms/step - loss: 1.5394 -
accuracy: 0.3885 - val_loss: 1.5664 - val_accuracy: 0.3790
Epoch 12/20
628/628 [=====] - 2s 3ms/step - loss: 1.5223 -
accuracy: 0.3957 - val_loss: 1.5696 - val_accuracy: 0.3929
Epoch 13/20
628/628 [=====] - 2s 3ms/step - loss: 1.5058 -
accuracy: 0.4062 - val_loss: 1.5577 - val_accuracy: 0.3878
Epoch 14/20
628/628 [=====] - 2s 3ms/step - loss: 1.4931 -
accuracy: 0.4085 - val_loss: 1.5376 - val_accuracy: 0.3990
Epoch 15/20
628/628 [=====] - 2s 3ms/step - loss: 1.4742 -
accuracy: 0.4136 - val_loss: 1.5383 - val_accuracy: 0.3950
Epoch 16/20
628/628 [=====] - 2s 3ms/step - loss: 1.4651 -
accuracy: 0.4182 - val_loss: 1.5363 - val_accuracy: 0.3990
Epoch 17/20
628/628 [=====] - 2s 3ms/step - loss: 1.4434 -
accuracy: 0.4314 - val_loss: 1.5395 - val_accuracy: 0.3890
Epoch 18/20
628/628 [=====] - 2s 3ms/step - loss: 1.4410 -
accuracy: 0.4343 - val_loss: 1.5507 - val_accuracy: 0.3946
Epoch 19/20
628/628 [=====] - 2s 3ms/step - loss: 1.4220 -
accuracy: 0.4378 - val_loss: 1.5614 - val_accuracy: 0.3943
Epoch 20/20
628/628 [=====] - 2s 3ms/step - loss: 1.4018 -
accuracy: 0.4449 - val_loss: 1.5286 - val_accuracy: 0.4129


```

```

[44]: plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.title('model loss')
      plt.ylabel('loss')


```

```
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



```
[45]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model acc')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```






```
[43]: model.evaluate(test_x, test_y_cat)
```

```
135/135 [=====] - 0s 2ms/step - loss: 1.5302 -
accuracy: 0.4065
```

```
[43]: [1.5301593542099, 0.406547486782074]
```

```
[50]: cm = confusion_matrix(test_y, np.apply_along_axis(np.argmax, 1, model.
    ↪predict(test_x)))
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=nb.classes_, yticklabels=nb.
    ↪classes_)
plt.title('Confusion Matrix of ANN')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show(block=False)
```



## 2.7 Convolutional Neural Network

```
[51]: train_x_cnn = train_x.reshape(-1, 48, 48, 1)
      val_x_cnn = val_x.reshape(-1, 48, 48, 1)
      test_x_cnn = test_x.reshape(-1, 48, 48, 1)
```

```
[52]: model2 = Sequential()
      model2.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(48, 48, 1)))
      model2.add(layers.MaxPooling2D((2, 2)))
      model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```

model2.add(layers.Dropout(0.5))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Dropout(0.5))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.Flatten())
model2.add(layers.Dropout(0.2))
model2.add(layers.Dense(64, activation='relu'))
model2.add(layers.Dropout(0.5))
model2.add(layers.Dense(7))

```

```

[53]: model2.compile(optimizer='adam',
                    loss=tf.keras.losses.
                        ↪SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])
model2.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 46, 46, 32)	320
max_pooling2d (MaxPooling2D)	(None, 23, 23, 32)	0
conv2d_1 (Conv2D)	(None, 21, 21, 64)	18496
dropout_2 (Dropout)	(None, 21, 21, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_3 (Dropout)	(None, 10, 10, 64)	0
conv2d_2 (Conv2D)	(None, 8, 8, 64)	36928
flatten (Flatten)	(None, 4096)	0
dropout_4 (Dropout)	(None, 4096)	0
dense_3 (Dense)	(None, 64)	262208
dropout_5 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 7)	455

Total params: 318,407  
Trainable params: 318,407  
Non-trainable params: 0

```
[54]: history2 = model2.fit(train_x_cnn, train_y, epochs = 20, validation_data =  
    ↪(val_x_cnn, val_y))
```

Epoch 1/20  
628/628 [=====] - 12s 18ms/step - loss: 1.7455 -  
accuracy: 0.2952 - val\_loss: 1.6535 - val\_accuracy: 0.3994  
Epoch 2/20  
628/628 [=====] - 10s 16ms/step - loss: 1.6020 -  
accuracy: 0.3716 - val\_loss: 1.5413 - val\_accuracy: 0.4371  
Epoch 3/20  
628/628 [=====] - 10s 17ms/step - loss: 1.5322 -  
accuracy: 0.4031 - val\_loss: 1.4873 - val\_accuracy: 0.4603  
Epoch 4/20  
628/628 [=====] - 11s 17ms/step - loss: 1.4787 -  
accuracy: 0.4273 - val\_loss: 1.4146 - val\_accuracy: 0.4807  
Epoch 5/20  
628/628 [=====] - 10s 16ms/step - loss: 1.4423 -  
accuracy: 0.4407 - val\_loss: 1.3823 - val\_accuracy: 0.4868  
Epoch 6/20  
628/628 [=====] - 10s 16ms/step - loss: 1.4111 -  
accuracy: 0.4538 - val\_loss: 1.3680 - val\_accuracy: 0.4998  
Epoch 7/20  
628/628 [=====] - 10s 16ms/step - loss: 1.3862 -  
accuracy: 0.4660 - val\_loss: 1.3246 - val\_accuracy: 0.5109  
Epoch 8/20  
628/628 [=====] - 10s 17ms/step - loss: 1.3646 -  
accuracy: 0.4727 - val\_loss: 1.3322 - val\_accuracy: 0.5193  
Epoch 9/20  
628/628 [=====] - 10s 16ms/step - loss: 1.3412 -  
accuracy: 0.4868 - val\_loss: 1.3130 - val\_accuracy: 0.5128  
Epoch 10/20  
628/628 [=====] - 10s 16ms/step - loss: 1.3273 -  
accuracy: 0.4917 - val\_loss: 1.2800 - val\_accuracy: 0.5260  
Epoch 11/20  
628/628 [=====] - 10s 16ms/step - loss: 1.3171 -  
accuracy: 0.4952 - val\_loss: 1.2878 - val\_accuracy: 0.5251  
Epoch 12/20  
628/628 [=====] - 10s 16ms/step - loss: 1.2988 -  
accuracy: 0.4991 - val\_loss: 1.2500 - val\_accuracy: 0.5307  
Epoch 13/20  
628/628 [=====] - 10s 16ms/step - loss: 1.2877 -  
accuracy: 0.5042 - val\_loss: 1.2738 - val\_accuracy: 0.5381  
Epoch 14/20

```


628/628 [=====] - 10s 16ms/step - loss: 1.2815 -
accuracy: 0.5094 - val_loss: 1.2709 - val_accuracy: 0.5390
Epoch 15/20
628/628 [=====] - 10s 16ms/step - loss: 1.2663 -
accuracy: 0.5157 - val_loss: 1.2694 - val_accuracy: 0.5437
Epoch 16/20
628/628 [=====] - 10s 16ms/step - loss: 1.2596 -
accuracy: 0.5146 - val_loss: 1.2666 - val_accuracy: 0.5439
Epoch 17/20
628/628 [=====] - 10s 16ms/step - loss: 1.2487 -
accuracy: 0.5222 - val_loss: 1.2444 - val_accuracy: 0.5444
Epoch 18/20
628/628 [=====] - 10s 16ms/step - loss: 1.2448 -
accuracy: 0.5208 - val_loss: 1.2193 - val_accuracy: 0.5462
Epoch 19/20
628/628 [=====] - 10s 16ms/step - loss: 1.2365 -
accuracy: 0.5246 - val_loss: 1.2314 - val_accuracy: 0.5467
Epoch 20/20
628/628 [=====] - 10s 17ms/step - loss: 1.2251 -
accuracy: 0.5249 - val_loss: 1.2304 - val_accuracy: 0.5476

```


```

[55]: plt.plot(history2.history['loss'])
plt.plot(history2.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()

```



```
[56]: plt.plot(history2.history['accuracy'])
plt.plot(history2.history['val_accuracy'])
plt.title('model acc')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```




```
[57]: model2.evaluate(test_x_cnn, test_y)
```

```
135/135 [=====] - 1s 5ms/step - loss: 1.2294 -
accuracy: 0.5482
```

```
[57]: [1.229351282119751, 0.5481773614883423]
```

```
[81]: cm = confusion_matrix(test_y, np.apply_along_axis(np.argmax, 1, model2.
    ↪predict(test_x_cnn)))
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=nb.classes_, yticklabels=nb.
    ↪classes_)
plt.title('Confusion Matrix of CNN')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show(block=False)
```



```
[83]: model.save('saved_model/ann')  
      model2.save('saved_model/cnn')
```

```
INFO:tensorflow:Assets written to: saved_model/ann/assets  
INFO:tensorflow:Assets written to: saved_model/cnn/assets
```